

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
sns.set(color_codes = True)
sns.set(style="whitegrid")
sns.set(rc={'figure.figsize':(7,4)})
sns.set_palette("Set3")
import warnings
warnings.filterwarnings('ignore')
warnings.filterwarnings(action='ignore', category=DeprecationWarning)
pd.set_option('display.max_columns', None)
pd.set_option('display.float_format', lambda x: '%.2f' % x)
```

```
In [2]: df = pd.read_csv("high_churn_list_model.csv") # data = pd.read_csv("census.csv")
```

```
In [3]: df.describe(include = ['O']).transpose()[1:]
```

Out[3]:

|                             | count | unique | top             | freq |
|-----------------------------|-------|--------|-----------------|------|
| <b>City</b>                 | 2095  | 876    | Los Angeles     | 93   |
| <b>Offer</b>                | 2095  | 6      | None            | 1152 |
| <b>MultipleLines</b>        | 2095  | 3      | No              | 1005 |
| <b>InternetType</b>         | 2095  | 4      | Fiber Optic     | 933  |
| <b>OnlineSecurity</b>       | 2095  | 3      | No              | 1065 |
| <b>OnlineBackup</b>         | 2095  | 3      | No              | 920  |
| <b>DeviceProtectionPlan</b> | 2095  | 3      | No              | 956  |
| <b>PremiumTechSupport</b>   | 2095  | 3      | No              | 1064 |
| <b>StreamingTV</b>          | 2095  | 3      | No              | 865  |
| <b>StreamingMovies</b>      | 2095  | 3      | No              | 842  |
| <b>StreamingMusic</b>       | 2095  | 3      | No              | 929  |
| <b>UnlimitedData</b>        | 2095  | 3      | Yes             | 1406 |
| <b>Contract</b>             | 2095  | 3      | Month-to-Month  | 1096 |
| <b>PaymentMethod</b>        | 2095  | 3      | Bank Withdrawal | 1171 |
| <b>ChurnCategory</b>        | 2095  | 6      | Unknown         | 1537 |
| <b>ChurnReason</b>          | 2095  | 21     | Unknown         | 1537 |

```
In [4]: df.describe(exclude = ['O']).transpose()[:-3]
```

Out[4]:

|                                      | count   | mean     | std      | min      | 25%      | 50%      | 75%    |
|--------------------------------------|---------|----------|----------|----------|----------|----------|--------|
| <b>Gender</b>                        | 2095.00 | 0.50     | 0.50     | 0.00     | 0.00     | 1.00     | 1.     |
| <b>Age</b>                           | 2095.00 | 46.68    | 16.80    | 19.00    | 32.00    | 46.00    | 60.    |
| <b>Married</b>                       | 2095.00 | 0.48     | 0.50     | 0.00     | 0.00     | 0.00     | 1.     |
| <b>NumberOfDependents</b>            | 2095.00 | 0.46     | 0.98     | 0.00     | 0.00     | 0.00     | 0.     |
| <b>ZipCode</b>                       | 2095.00 | 93521.26 | 1860.72  | 90001.00 | 92109.00 | 93550.00 | 95359. |
| <b>Population</b>                    | 2095.00 | 21501.67 | 20232.50 | 11.00    | 2347.50  | 16717.00 | 35109. |
| <b>NumberOfReferrals</b>             | 2095.00 | 1.95     | 3.00     | 0.00     | 0.00     | 0.00     | 3.     |
| <b>TenureinMonths</b>                | 2095.00 | 31.95    | 24.33    | 1.00     | 8.00     | 28.00    | 55.    |
| <b>PhoneService</b>                  | 2095.00 | 0.91     | 0.29     | 0.00     | 1.00     | 1.00     | 1.     |
| <b>AvgMonthlyLongDistanceCharges</b> | 2095.00 | 22.78    | 15.45    | 0.00     | 9.27     | 22.30    | 36.    |
| <b>InternetService</b>               | 2095.00 | 0.79     | 0.41     | 0.00     | 1.00     | 1.00     | 1.     |
| <b>AvgMonthlyGBDownload</b>          | 2095.00 | 20.72    | 20.59    | 0.00     | 4.00     | 17.00    | 27.    |
| <b>PaperlessBilling</b>              | 2095.00 | 0.58     | 0.49     | 0.00     | 0.00     | 1.00     | 1.     |
| <b>MonthlyCharge</b>                 | 2095.00 | 63.68    | 30.97    | -10.00   | 33.52    | 70.35    | 89.    |
| <b>TotalCharges</b>                  | 2095.00 | 2242.08  | 2223.44  | 18.80    | 406.27   | 1387.45  | 3704.  |
| <b>TotalRefunds</b>                  | 2095.00 | 1.82     | 7.54     | 0.00     | 0.00     | 0.00     | 0.     |
| <b>TotalExtraDataCharges</b>         | 2095.00 | 7.26     | 25.99    | 0.00     | 0.00     | 0.00     | 0.     |
| <b>TotalLongDistanceCharges</b>      | 2095.00 | 734.42   | 833.37   | 0.00     | 74.54    | 378.00   | 1187.  |
| <b>TotalRevenue</b>                  | 2095.00 | 2981.95  | 2813.69  | 21.36    | 589.41   | 2047.16  | 4733.  |

## Determining how many "churned"

In [5]: `df['Churn'].value_counts()`

Out[5]:

|   |      |
|---|------|
| 0 | 1537 |
| 1 | 558  |

Name: Churn, dtype: int64

In [6]: `df.groupby(['Married']).agg({'Churn': 'mean'}).reset_index().sort_values(by='Churn')`

Out[6]:

|   | Married | Churn |
|---|---------|-------|
| 0 | 0       | 0.33  |
| 1 | 1       | 0.20  |

In [7]: `df.groupby(['Gender']).agg({'Churn': 'mean'}).reset_index().sort_values(by='Churn',`

Out[7]:

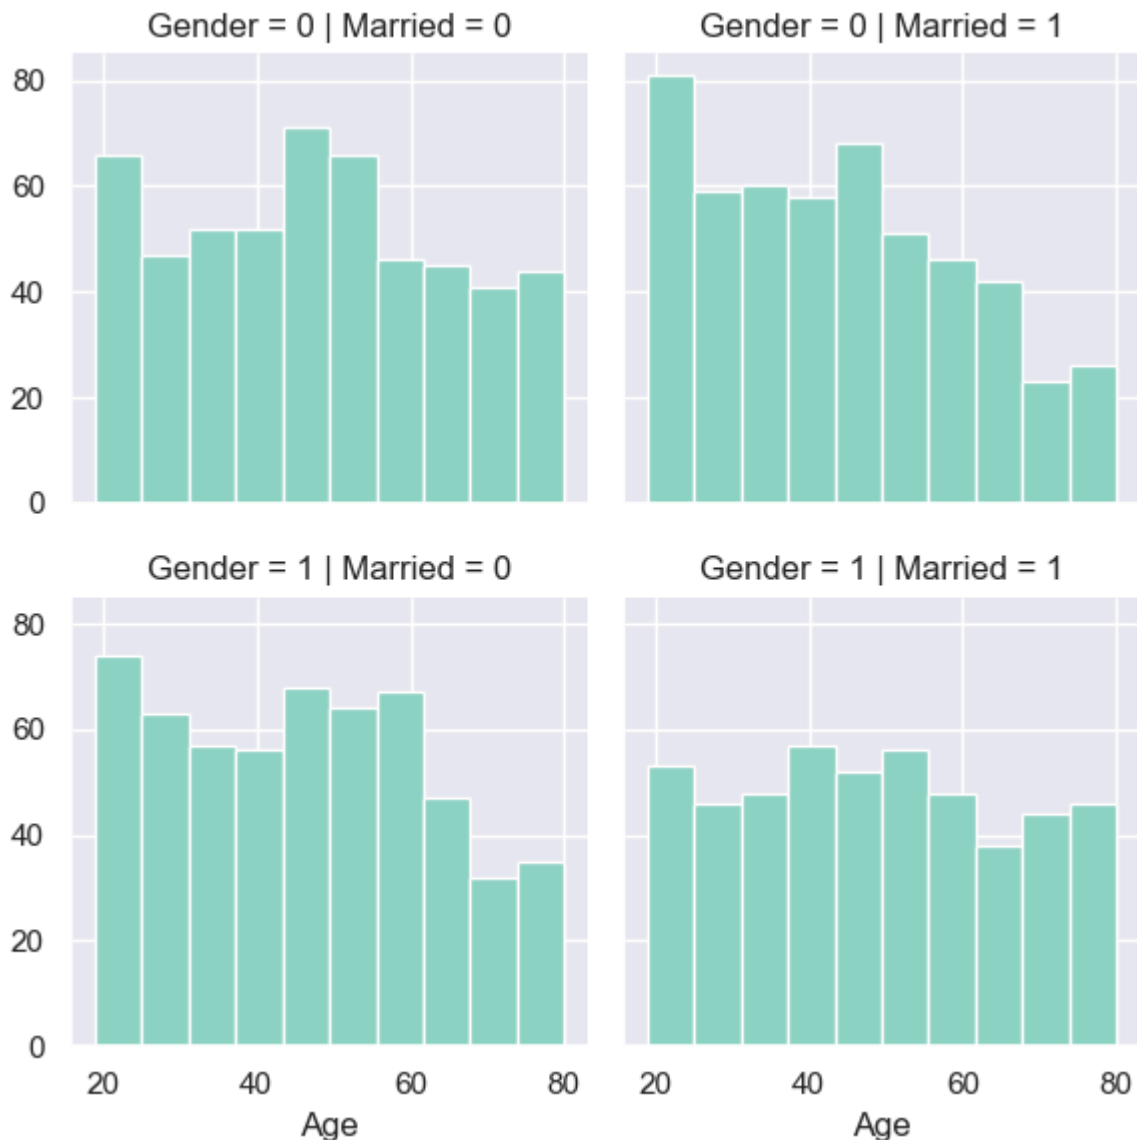
|   | Gender | Churn |
|---|--------|-------|
| 0 | 0      | 0.27  |
| 1 | 1      | 0.26  |

```
In [8]: df.groupby(['Gender', 'Married']).agg({'Churn': 'mean'}).reset_index().sort_values(b
```

```
Out[8]:
```

|   | Gender | Married | Churn |
|---|--------|---------|-------|
| 0 | 0      | 0       | 0.35  |
| 2 | 1      | 0       | 0.31  |
| 3 | 1      | 1       | 0.20  |
| 1 | 0      | 1       | 0.19  |

```
In [9]: g = sns.FacetGrid(df, col = "Married", row = 'Gender')
g = g.map(plt.hist, "Age")
```



```
In [10]: fig, [ax0, ax1, ax2] = plt.subplots(1,3, figsize = (14,4))

ax0.hist(df['TotalRevenue'])
ax0.set_xlabel('Total Revenue Distribution')
ax0.axvline(df['TotalRevenue'].mean(), color = "black")

ax1.hist(df.TotalCharges)
ax1.set_xlabel('Total Charges Distribution')
ax1.axvline(df["TotalCharges"].mean(), color = "black");

ax2.hist(df.TenureinMonths)
```

```
ax2.set_xlabel('Tenure in Month distribution')
ax2.axvline(df['TenureinMonths'].mean(), color = "black");

print("Black lines are means")
```

Black lines are means

```
In [11]: fig, [ax0, ax1, ax2] = plt.subplots(1,3, figsize = (14,4))

ax0.hist(df['Age'])
ax0.set_xlabel('Age Distribution')
ax0.axvline(df['Age'].mean(), color = "black")

ax1.hist(df['TotalLongDistanceCharges'])
ax1.set_xlabel('Total Long Distance Charges Distribution')
ax1.axvline(df["TotalLongDistanceCharges"].mean(), color = "black");

ax2.hist(df['Population'])
ax2.set_xlabel('Population Distribution')
ax2.axvline(df['Population'].mean(), color = "black");

print("Black lines are means")
```

Black lines are means

```
In [12]: sns.countplot(y = df['TotalExtraDataCharges'])
```

```
Out[12]: <AxesSubplot:xlabel='count', ylabel='TotalExtraDataCharges'>
```

```
In [13]: sns.violinplot(x=df['TotalRefunds'])
```

```
Out[13]: <AxesSubplot:xlabel='TotalRefunds', ylabel='TotalExtraDataCharges'>
```

```
In [14]: sns.boxplot(x=df['MonthlyCharge'])
```

```
Out[14]: <AxesSubplot:xlabel='MonthlyCharge', ylabel='TotalExtraDataCharges'>
```

```
In [15]: sns.countplot(df['PaymentMethod'])
```

```
Out[15]: <AxesSubplot:xlabel='PaymentMethod', ylabel='count'>
```

```
In [16]: sns.countplot(y=df['ChurnCategory'],order = df['ChurnCategory'].value_counts().index)
```

```
Out[16]: <AxesSubplot:xlabel='count', ylabel='ChurnCategory'>
```

```
In [17]: plt.figure(figsize=(10,10))
sns.countplot(y=df['ChurnReason'],order = df['ChurnReason'].value_counts().index)
```

```
Out[17]: <AxesSubplot:xlabel='count', ylabel='ChurnReason'>
```

```
In [18]: sns.countplot(df['PhoneService'])
```

```
Out[18]: <AxesSubplot:xlabel='PhoneService', ylabel='count'>
```

```
In [19]: sns.countplot(df['PaperlessBilling'],hue=df['Churn'])
```

```
Out[19]: <AxesSubplot:xlabel='PaperlessBilling', ylabel='count'>
```

```
In [20]: sns.countplot(df['Contract'],hue = df['Churn'],order = df['Contract'].value_counts(
```

```
Out[20]: <AxesSubplot:xlabel='Contract', ylabel='count'>
```

```
In [21]: sns.countplot(df['Offer'], hue = df['Churn'], order = df['Offer'].value_counts().index)
```

```
Out[21]: <AxesSubplot:xlabel='Offer', ylabel='count'>
```

```
In [22]: sns.countplot(df['InternetType'], order = df['InternetType'].value_counts().index)
```

```
Out[22]: <AxesSubplot:xlabel='InternetType', ylabel='count'>
```

## Exploratory data analysis

```
In [23]: quant_df1 = df[[ #nominal
    'Churn',
    'Age',
    'NumberOfDependents',
    'Population',
    'NumberOfReferrals',
    'TenureinMonths',
    'AvgMonthlyLongDistanceCharges',
    'AvgMonthlyGBDownload',
    'MonthlyCharge',
    'TotalCharges',
    'TotalRefunds',
    'TotalExtraDataCharges',
    'TotalLongDistanceCharges',
    'TotalRevenue'
]].copy()

# categorical columns

quant_df2 = df[[ #binary
    'Churn',
    'Gender',
    'Married',
    'PhoneService',
    'InternetService',
    'PaperlessBilling'
]].copy()

quant_df3 = df[[ #categories (one-hot-encode)
    'Churn',
    'Offer',
    'MultipleLines',
    'InternetType',
    'OnlineSecurity',
    'OnlineBackup',
    'DeviceProtectionPlan',
    'PremiumTechSupport',
    'StreamingTV',
    'StreamingMovies',
    'StreamingMusic',
    'UnlimitedData',
    'Contract',
    'PaymentMethod'
]].copy()
```

```
In [24]: quant_df1.corr()
```

Out[24]:

|                               | Churn | Age   | NumberofDependents | Population | NumberofReferr |
|-------------------------------|-------|-------|--------------------|------------|----------------|
| Churn                         | 1.00  | 0.13  | -0.21              | 0.08       | -0             |
| Age                           | 0.13  | 1.00  | -0.09              | -0.02      | -0             |
| NumberofDependents            | -0.21 | -0.09 | 1.00               | -0.02      | 0              |
| Population                    | 0.08  | -0.02 | -0.02              | 1.00       | -0             |
| NumberofReferrals             | -0.30 | -0.04 | 0.28               | -0.04      | 1              |
| TenureinMonths                | -0.37 | -0.01 | 0.14               | -0.05      | 0              |
| AvgMonthlyLongDistanceCharges | 0.02  | -0.00 | -0.01              | -0.04      | 0              |
| AvgMonthlyGBDownload          | 0.02  | -0.36 | 0.17               | 0.01       | 0              |
| MonthlyCharge                 | 0.16  | 0.14  | -0.10              | -0.00      | 0              |
| TotalCharges                  | -0.22 | 0.05  | 0.04               | -0.04      | 0              |
| TotalRefunds                  | -0.02 | 0.00  | -0.01              | 0.05       | 0              |
| TotalExtraDataCharges         | 0.02  | 0.05  | -0.01              | 0.01       | -0             |
| TotalLongDistanceCharges      | -0.22 | 0.01  | 0.09               | -0.05      | 0              |
| TotalRevenue                  | -0.24 | 0.04  | 0.06               | -0.05      | 0              |

```
In [25]: plt.figure(figsize=(10,10))
cmap = sns.diverging_palette(250, 10, as_cmap=True)
sns.heatmap(quant_df1.corr(), cmap = cmap, annot = True);
```

```
In [26]: quant_df1.corr()['Churn'][1:]
```

```
Out[26]: Age                0.13
NumberofDependents        -0.21
Population                 0.08
NumberofReferrals         -0.30
TenureinMonths            -0.37
AvgMonthlyLongDistanceCharges  0.02
AvgMonthlyGBDownload       0.02
MonthlyCharge             0.16
TotalCharges              -0.22
TotalRefunds              -0.02
TotalExtraDataCharges      0.02
TotalLongDistanceCharges  -0.22
TotalRevenue              -0.24
Name: Churn, dtype: float64
```

```
In [27]: print(quant_df1.columns)

Index(['Churn', 'Age', 'NumberofDependents', 'Population', 'NumberofReferrals',
      'TenureinMonths', 'AvgMonthlyLongDistanceCharges',
      'AvgMonthlyGBDownload', 'MonthlyCharge', 'TotalCharges', 'TotalRefunds',
      'TotalExtraDataCharges', 'TotalLongDistanceCharges', 'TotalRevenue'],
      dtype='object')
```

Let's check our confidense about this statment with logistic regression model:

```
In [28]: quant_df1['intercept'] = 1
log_mod = sm.Logit(quant_df1['Churn'], quant_df1[['intercept', 'Age', 'NumberofDeper
      'TenureinMonths', 'AvgMonthlyLongDistanceCharges', 'AvgMonthlyGBDownload', 'M
```

```
'TotalExtraDataCharges','TotalRevenue']]).fit()
log_mod.summary()
```

Optimization terminated successfully.  
Current function value: 0.419326  
Iterations 7

Out[28]:

Logit Regression Results

|                         |                  |                          |            |
|-------------------------|------------------|--------------------------|------------|
| <b>Dep. Variable:</b>   | Churn            | <b>No. Observations:</b> | 2095       |
| <b>Model:</b>           | Logit            | <b>Df Residuals:</b>     | 2082       |
| <b>Method:</b>          | MLE              | <b>Df Model:</b>         | 12         |
| <b>Date:</b>            | Sat, 17 Aug 2024 | <b>Pseudo R-squ.:</b>    | 0.2765     |
| <b>Time:</b>            | 07:55:13         | <b>Log-Likelihood:</b>   | -878.49    |
| <b>converged:</b>       | True             | <b>LL-Null:</b>          | -1214.2    |
| <b>Covariance Type:</b> | nonrobust        | <b>LLR p-value:</b>      | 5.482e-136 |

|                                      | coef      | std err  | z      | P> z  | [0.025   | 0.975]   |
|--------------------------------------|-----------|----------|--------|-------|----------|----------|
| <b>intercept</b>                     | -1.5502   | 0.304    | -5.092 | 0.000 | -2.147   | -0.954   |
| <b>Age</b>                           | 0.0152    | 0.004    | 3.779  | 0.000 | 0.007    | 0.023    |
| <b>NumberofDependents</b>            | -0.4782   | 0.103    | -4.657 | 0.000 | -0.679   | -0.277   |
| <b>Population</b>                    | 7.153e-06 | 2.87e-06 | 2.495  | 0.013 | 1.53e-06 | 1.28e-05 |
| <b>NumberofReferrals</b>             | -0.2881   | 0.038    | -7.628 | 0.000 | -0.362   | -0.214   |
| <b>TenureinMonths</b>                | -0.0677   | 0.010    | -6.472 | 0.000 | -0.088   | -0.047   |
| <b>AvgMonthlyLongDistanceCharges</b> | -0.0113   | 0.006    | -1.995 | 0.046 | -0.022   | -0.000   |
| <b>AvgMonthlyGBDownload</b>          | 0.0055    | 0.004    | 1.487  | 0.137 | -0.002   | 0.013    |
| <b>MonthlyCharge</b>                 | 0.0218    | 0.003    | 6.772  | 0.000 | 0.015    | 0.028    |
| <b>TotalCharges</b>                  | -0.0005   | 0.000    | -1.876 | 0.061 | -0.001   | 2.03e-05 |
| <b>TotalRefunds</b>                  | -0.0061   | 0.008    | -0.757 | 0.449 | -0.022   | 0.010    |
| <b>TotalExtraDataCharges</b>         | 0.0006    | 0.002    | 0.280  | 0.779 | -0.004   | 0.005    |
| <b>TotalRevenue</b>                  | 0.0006    | 0.000    | 2.849  | 0.004 | 0.000    | 0.001    |

```
In [29]: # p values
log_mod.pvalues[:].plot.bar()
plt.axhline(y = 0.05);
```

```
In [30]: #coefficient
log_mod.params[:].plot.bar()
plt.axhline(y = 0.05);
```

```
In [31]: quant_df_main = {}
for i in log_mod.params[:].to_dict().keys():
    if log_mod.pvalues[i] < 0.05:
        quant_df_main[i] = log_mod.params[i]
    else:
        continue
quant_df_main
sorted(quant_df_main.items(), key=lambda x: x[1]) #sorting by highest to lowest val
```

```
Out[31]: [('intercept', -1.5501574689605269),
          ('NumberofDependents', -0.4781870547717154),
          ('NumberofReferrals', -0.2880830533840693),
          ('TenureinMonths', -0.06765171718218492),
          ('AvgMonthlyLongDistanceCharges', -0.01125946374212831),
          ('Population', 7.153223010203917e-06),
          ('TotalRevenue', 0.0005601071073698724),
          ('Age', 0.01521105759961699),
          ('MonthlyCharge', 0.021755369254524072)]
```

In [ ]:

## Compute the odds

```
In [32]: quant_df_main_odds = {k : np.exp(v) for k, v in quant_df_main.items()}
         sorted(quant_df_main_odds.items(), key=lambda x: x[1]) #sorting by highest to lowest
```

```
Out[32]: [('intercept', 0.21221455399029313),
          ('NumberofDependents', 0.6199062297179322),
          ('NumberofReferrals', 0.749699324587366),
          ('TenureinMonths', 0.9345859171174311),
          ('AvgMonthlyLongDistanceCharges', 0.9888036867838459),
          ('Population', 1.0000071532485946),
          ('TotalRevenue', 1.000560263996646),
          ('Age', 1.0153273345533307),
          ('MonthlyCharge', 1.0219937427965726)]
```

```
In [33]: quant_df1['intercept'] = 1
         log_mod2 = sm.Logit(quant_df1['Churn'], quant_df1[['intercept', 'Age', 'Population',
         log_mod2.summary()
```

Optimization terminated successfully.  
Current function value: 0.475116  
Iterations 6

Out[33]: Logit Regression Results

|                         |                  |                          |           |                 |               |               |
|-------------------------|------------------|--------------------------|-----------|-----------------|---------------|---------------|
| <b>Dep. Variable:</b>   | Churn            | <b>No. Observations:</b> | 2095      |                 |               |               |
| <b>Model:</b>           | Logit            | <b>Df Residuals:</b>     | 2090      |                 |               |               |
| <b>Method:</b>          | MLE              | <b>Df Model:</b>         | 4         |                 |               |               |
| <b>Date:</b>            | Sat, 17 Aug 2024 | <b>Pseudo R-squ.:</b>    | 0.1803    |                 |               |               |
| <b>Time:</b>            | 07:55:13         | <b>Log-Likelihood:</b>   | -995.37   |                 |               |               |
| <b>converged:</b>       | True             | <b>LL-Null:</b>          | -1214.2   |                 |               |               |
| <b>Covariance Type:</b> | nonrobust        | <b>LLR p-value:</b>      | 1.925e-93 |                 |               |               |
|                         | <b>coef</b>      | <b>std err</b>           | <b>z</b>  | <b>P&gt; z </b> | <b>[0.025</b> | <b>0.975]</b> |
| <b>intercept</b>        | -2.9004          | 0.216                    | -13.451   | 0.000           | -3.323        | -2.478        |
| <b>Age</b>              | 0.0140           | 0.003                    | 4.211     | 0.000           | 0.007         | 0.020         |
| <b>Population</b>       | 6.696e-06        | 2.69e-06                 | 2.491     | 0.013           | 1.43e-06      | 1.2e-05       |
| <b>MonthlyCharge</b>    | 0.0342           | 0.002                    | 14.282    | 0.000           | 0.029         | 0.039         |
| <b>TotalRevenue</b>     | -0.0005          | 2.88e-05                 | -16.004   | 0.000           | -0.001        | -0.000        |

## Log Mod 2 (Binary)



In [34]: `quant_df2.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2095 entries, 0 to 2094
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                2095 non-null   int64
1   Married               2095 non-null   int64
2   PhoneService          2095 non-null   int64
3   InternetService       2095 non-null   int64
4   PaperlessBilling      2095 non-null   int64
dtypes: int64(5)
memory usage: 82.0 KB
```

In [35]: `quant_df1['intercept'] = 1`  
`log_mod2 = sm.Logit(quant_df1['Churn'], quant_df2).fit()`  
`log_mod2.summary()`

Optimization terminated successfully.  
 Current function value: 0.562800  
 Iterations 5

Out[35]:

Logit Regression Results

|                         |                  |                          |           |
|-------------------------|------------------|--------------------------|-----------|
| <b>Dep. Variable:</b>   | Churn            | <b>No. Observations:</b> | 2095      |
| <b>Model:</b>           | Logit            | <b>Df Residuals:</b>     | 2090      |
| <b>Method:</b>          | MLE              | <b>Df Model:</b>         | 4         |
| <b>Date:</b>            | Sat, 17 Aug 2024 | <b>Pseudo R-squ.:</b>    | 0.02897   |
| <b>Time:</b>            | 07:55:13         | <b>Log-Likelihood:</b>   | -1179.1   |
| <b>converged:</b>       | True             | <b>LL-Null:</b>          | -1214.2   |
| <b>Covariance Type:</b> | nonrobust        | <b>LLR p-value:</b>      | 1.902e-14 |

|                         | coef    | std err | z      | P> z  | [0.025 | 0.975] |
|-------------------------|---------|---------|--------|-------|--------|--------|
| <b>Gender</b>           | -0.3192 | 0.098   | -3.271 | 0.001 | -0.510 | -0.128 |
| <b>Married</b>          | -0.9272 | 0.101   | -9.148 | 0.000 | -1.126 | -0.729 |
| <b>PhoneService</b>     | -0.9664 | 0.106   | -9.117 | 0.000 | -1.174 | -0.759 |
| <b>InternetService</b>  | 0.2870  | 0.115   | 2.490  | 0.013 | 0.061  | 0.513  |
| <b>PaperlessBilling</b> | 0.4793  | 0.110   | 4.355  | 0.000 | 0.264  | 0.695  |

In [36]: `quant_df_main2 = {}`  
`for i in log_mod2.params[:].to_dict().keys():`  
 `if log_mod2.pvalues[i] < 0.05:`  
 `quant_df_main2[i] = log_mod2.params[i].round(4)`  
 `else:`  
 `continue`  
`quant_df_main2`  
`sorted(quant_df_main2.items(), key=lambda x: x[1]) #sorting by highe`Out[36]: `[('PhoneService', -0.9664),`  
 `('Married', -0.9272),`  
 `('Gender', -0.3192),`  
 `('InternetService', 0.287),`  
 `('PaperlessBilling', 0.4793)]`In [37]: `quant_df_main_odds2 = {k : np.exp(v) for k, v in quant_df_main2.items()}`  
`sorted(quant_df_main_odds2.items(), key=lambda x: x[1])`

```
Out[37]: [('PhoneService', 0.3804501964490625),
          ('Married', 0.3956600088557378),
          ('Gender', 0.7267301887330189),
          ('InternetService', 1.3324242134756759),
          ('PaperlessBilling', 1.6149435459572175)]
```

## Log Mod 3 (Category)

```
In [38]: quant_df3 = df[[
          'Churn',
          'Offer',
          'InternetType',
          'Contract',
          'PaymentMethod',
          ]].copy()
```

```
In [39]: quant_df3.head()
```

```
Out[39]:
```

|   | Churn | Offer   | InternetType | Contract       | PaymentMethod   |
|---|-------|---------|--------------|----------------|-----------------|
| 0 | 1     | Offer E | Fiber Optic  | Month-to-Month | Bank Withdrawal |
| 1 | 1     | None    | Fiber Optic  | Month-to-Month | Bank Withdrawal |
| 2 | 1     | None    | Fiber Optic  | Month-to-Month | Bank Withdrawal |
| 3 | 1     | None    | DSL          | Month-to-Month | Bank Withdrawal |
| 4 | 1     | Offer D | Fiber Optic  | Month-to-Month | Credit Card     |

```
In [40]: quant_df3 = pd.get_dummies(quant_df3)
```

```
In [41]: quant_df3.describe().transpose()
```

Out[41]:

|                                      | count   | mean | std  | min  | 25%  | 50%  | 75%  | max  |
|--------------------------------------|---------|------|------|------|------|------|------|------|
| <b>Churn</b>                         | 2095.00 | 0.27 | 0.44 | 0.00 | 0.00 | 0.00 | 1.00 | 1.00 |
| <b>Offer_None</b>                    | 2095.00 | 0.55 | 0.50 | 0.00 | 0.00 | 1.00 | 1.00 | 1.00 |
| <b>Offer_Offer A</b>                 | 2095.00 | 0.07 | 0.25 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>Offer_Offer B</b>                 | 2095.00 | 0.12 | 0.32 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>Offer_Offer C</b>                 | 2095.00 | 0.06 | 0.23 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>Offer_Offer D</b>                 | 2095.00 | 0.09 | 0.28 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>Offer_Offer E</b>                 | 2095.00 | 0.12 | 0.32 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>InternetType_Cable</b>            | 2095.00 | 0.12 | 0.32 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>InternetType_DSL</b>              | 2095.00 | 0.22 | 0.42 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>InternetType_Fiber Optic</b>      | 2095.00 | 0.45 | 0.50 | 0.00 | 0.00 | 0.00 | 1.00 | 1.00 |
| <b>InternetType_Unknown</b>          | 2095.00 | 0.21 | 0.41 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>Contract_Month-to-Month</b>       | 2095.00 | 0.52 | 0.50 | 0.00 | 0.00 | 1.00 | 1.00 | 1.00 |
| <b>Contract_One Year</b>             | 2095.00 | 0.21 | 0.41 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |
| <b>Contract_Two Year</b>             | 2095.00 | 0.27 | 0.44 | 0.00 | 0.00 | 0.00 | 1.00 | 1.00 |
| <b>PaymentMethod_Bank Withdrawal</b> | 2095.00 | 0.56 | 0.50 | 0.00 | 0.00 | 1.00 | 1.00 | 1.00 |
| <b>PaymentMethod_Credit Card</b>     | 2095.00 | 0.40 | 0.49 | 0.00 | 0.00 | 0.00 | 1.00 | 1.00 |
| <b>PaymentMethod_Mailed Check</b>    | 2095.00 | 0.05 | 0.21 | 0.00 | 0.00 | 0.00 | 0.00 | 1.00 |

In [42]: quant\_df3.shape

Out[42]: (2095, 17)

```
In [43]: log_mod3 = sm.Logit(quant_df3['Churn'], quant_df3[['Offer_None', 'Offer_Offer A', 'Of
                                                    'Offer_Offer C', 'Offer_Offer E',
log_mod3.summary()
```

Optimization terminated successfully.  
 Current function value: 0.542199  
 Iterations 7

Out[43]:

Logit Regression Results

|                  |                  |           |                   |       |           |        |
|------------------|------------------|-----------|-------------------|-------|-----------|--------|
| Dep. Variable:   |                  | Churn     | No. Observations: |       | 2095      |        |
| Model:           |                  | Logit     | Df Residuals:     |       | 2089      |        |
| Method:          |                  | MLE       | Df Model:         |       | 5         |        |
| Date:            | Sat, 17 Aug 2024 |           | Pseudo R-squ.:    |       | 0.06452   |        |
| Time:            | 07:55:14         |           | Log-Likelihood:   |       | -1135.9   |        |
| converged:       |                  | True      | LL-Null:          |       | -1214.2   |        |
| Covariance Type: |                  | nonrobust | LLR p-value:      |       | 5.044e-32 |        |
|                  |                  |           |                   |       |           |        |
|                  | coef             | std err   | z                 | P> z  | [0.025    | 0.975] |
| Offer_None       | -0.9729          | 0.066     | -14.733           | 0.000 | -1.102    | -0.843 |
| Offer_Offer A    | -2.9886          | 0.387     | -7.715            | 0.000 | -3.748    | -2.229 |
| Offer_Offer B    | -2.0384          | 0.201     | -10.146           | 0.000 | -2.432    | -1.645 |
| Offer_Offer C    | -1.2040          | 0.219     | -5.487            | 0.000 | -1.634    | -0.774 |
| Offer_Offer E    | 0.0877           | 0.126     | 0.694             | 0.488 | -0.160    | 0.335  |
| Offer_Offer D    | -1.0282          | 0.166     | -6.177            | 0.000 | -1.354    | -0.702 |

In [44]:

```
log_mod3 = sm.Logit(quant_df3['Churn'], quant_df3[['Offer_None', 'Offer_Offer A', 'Offer_Offer B', 'Offer_Offer C', 'Offer_Offer E', 'Offer_Offer D']])
log_mod3.summary()
```

Optimization terminated successfully.  
Current function value: 0.542199  
Iterations 7

Out[44]:

Logit Regression Results

|                  |                  |           |                   |       |           |        |
|------------------|------------------|-----------|-------------------|-------|-----------|--------|
| Dep. Variable:   |                  | Churn     | No. Observations: |       | 2095      |        |
| Model:           |                  | Logit     | Df Residuals:     |       | 2089      |        |
| Method:          |                  | MLE       | Df Model:         |       | 5         |        |
| Date:            | Sat, 17 Aug 2024 |           | Pseudo R-squ.:    |       | 0.06452   |        |
| Time:            | 07:55:14         |           | Log-Likelihood:   |       | -1135.9   |        |
| converged:       |                  | True      | LL-Null:          |       | -1214.2   |        |
| Covariance Type: |                  | nonrobust | LLR p-value:      |       | 5.044e-32 |        |
|                  |                  |           |                   |       |           |        |
|                  | coef             | std err   | z                 | P> z  | [0.025    | 0.975] |
| Offer_None       | -0.9729          | 0.066     | -14.733           | 0.000 | -1.102    | -0.843 |
| Offer_Offer A    | -2.9886          | 0.387     | -7.715            | 0.000 | -3.748    | -2.229 |
| Offer_Offer B    | -2.0384          | 0.201     | -10.146           | 0.000 | -2.432    | -1.645 |
| Offer_Offer C    | -1.2040          | 0.219     | -5.487            | 0.000 | -1.634    | -0.774 |
| Offer_Offer E    | 0.0877           | 0.126     | 0.694             | 0.488 | -0.160    | 0.335  |
| Offer_Offer D    | -1.0282          | 0.166     | -6.177            | 0.000 | -1.354    | -0.702 |

In [45]:

```
quant_df_main = {}
for i in log_mod3.params[:,].to_dict().keys():
```

```

if log_mod3.pvalues[i] < 0.05:
    quant_df_main[i] = log_mod3.params[i].round(4)
else:
    continue
quant_df_main
sorted(quant_df_main.items(), key=lambda x: x[1]) #sorting by highest to lowest val

```

```

Out[45]: [('Offer_Offer A', -2.9886),
          ('Offer_Offer B', -2.0384),
          ('Offer_Offer C', -1.204),
          ('Offer_Offer D', -1.0282),
          ('Offer_None', -0.9729)]

```

```

In [46]: quant_df3['intercept'] = 1
log_mod3 = sm.Logit(quant_df3['Churn'], quant_df3[['Contract_Month-to-Month',
                                                  'Contract_One Year', 'Contract_Two Year']])
log_mod3.summary()

```

Optimization terminated successfully.  
 Current function value: 0.465441  
 Iterations 8

Out[46]: Logit Regression Results

|                         |                  |                          |            |
|-------------------------|------------------|--------------------------|------------|
| <b>Dep. Variable:</b>   | Churn            | <b>No. Observations:</b> | 2095       |
| <b>Model:</b>           | Logit            | <b>Df Residuals:</b>     | 2092       |
| <b>Method:</b>          | MLE              | <b>Df Model:</b>         | 2          |
| <b>Date:</b>            | Sat, 17 Aug 2024 | <b>Pseudo R-squ.:</b>    | 0.1970     |
| <b>Time:</b>            | 07:55:14         | <b>Log-Likelihood:</b>   | -975.10    |
| <b>converged:</b>       | True             | <b>LL-Null:</b>          | -1214.2    |
| <b>Covariance Type:</b> | nonrobust        | <b>LLR p-value:</b>      | 1.377e-104 |

|                                | coef    | std err | z       | P> z  | [0.025 | 0.975] |
|--------------------------------|---------|---------|---------|-------|--------|--------|
| <b>Contract_Month-to-Month</b> | -0.2088 | 0.061   | -3.437  | 0.001 | -0.328 | -0.090 |
| <b>Contract_One Year</b>       | -1.9148 | 0.143   | -13.377 | 0.000 | -2.195 | -1.634 |
| <b>Contract_Two Year</b>       | -3.9157 | 0.305   | -12.859 | 0.000 | -4.512 | -3.319 |

```

In [47]: quant_df3['intercept'] = 1
log_mod3 = sm.Logit(quant_df3['Churn'], quant_df3[['PaymentMethod_Bank Withdrawal',
                                                  'PaymentMethod_Mailed Check']]).fit()
log_mod3.summary()

```

Optimization terminated successfully.  
 Current function value: 0.554132  
 Iterations 6

Out[47]:

Logit Regression Results

|                  |                  |                   |           |
|------------------|------------------|-------------------|-----------|
| Dep. Variable:   | Churn            | No. Observations: | 2095      |
| Model:           | Logit            | Df Residuals:     | 2092      |
| Method:          | MLE              | Df Model:         | 2         |
| Date:            | Sat, 17 Aug 2024 | Pseudo R-squ.:    | 0.04393   |
| Time:            | 07:55:14         | Log-Likelihood:   | -1160.9   |
| converged:       | True             | LL-Null:          | -1214.2   |
| Covariance Type: | nonrobust        | LLR p-value:      | 6.826e-24 |

|                               | coef    | std err | z       | P> z  | [0.025 | 0.975] |
|-------------------------------|---------|---------|---------|-------|--------|--------|
| PaymentMethod_Bank Withdrawal | -0.6791 | 0.062   | -10.980 | 0.000 | -0.800 | -0.558 |
| PaymentMethod_Credit Card     | -1.7474 | 0.098   | -17.885 | 0.000 | -1.939 | -1.556 |
| PaymentMethod_Mailed Check    | -0.2754 | 0.207   | -1.330  | 0.184 | -0.681 | 0.131  |

In [ ]:

```
quant_df3['intercept'] = 1
log_mod3 = sm.Logit(quant_df1['Churn'], quant_df3[['InternetType_Cable','InternetTy
'InternetType_Fiber Optic','Inter
log_mod3.summary()
```

Optimization terminated successfully.  
Current function value: 0.536571  
Iterations 6

Out[48]:

Logit Regression Results

|                  |                  |                   |           |
|------------------|------------------|-------------------|-----------|
| Dep. Variable:   | Churn            | No. Observations: | 2095      |
| Model:           | Logit            | Df Residuals:     | 2091      |
| Method:          | MLE              | Df Model:         | 3         |
| Date:            | Sat, 17 Aug 2024 | Pseudo R-squ.:    | 0.07423   |
| Time:            | 07:55:14         | Log-Likelihood:   | -1124.1   |
| converged:       | True             | LL-Null:          | -1214.2   |
| Covariance Type: | nonrobust        | LLR p-value:      | 7.736e-39 |

|                          | coef    | std err | z       | P> z  | [0.025 | 0.975] |
|--------------------------|---------|---------|---------|-------|--------|--------|
| InternetType_Cable       | -1.2397 | 0.153   | -8.096  | 0.000 | -1.540 | -0.940 |
| InternetType_DSL         | -1.4655 | 0.118   | -12.391 | 0.000 | -1.697 | -1.234 |
| InternetType_Fiber Optic | -0.4064 | 0.067   | -6.080  | 0.000 | -0.537 | -0.275 |
| InternetType_Unknown     | -2.2687 | 0.162   | -13.997 | 0.000 | -2.586 | -1.951 |

In [ ]:

In [ ]: